



1. Examine this YAML file and SQL query, what will this SQL query compile to?

Examine this YAML file:

```
version: 2
sources:
  - name: event_tickets
    database: raw
    schema: ticket_tailor
    tables:
      - name: orders
      - name: events
      - name: tickets
        identifier: issued_tickets
```

Examine this SQL query:

```
select * from {{ source('event_tickets', 'tickets') }}
```

What will this SQL query compile to?

select * from raw. ☒ 'event_tickets' ☒ 'tickets'

☐ 'ticket_tailor' ☐ 'issued_tickets'

2. Which command will build only the models that were not built on the recent run?
- A: dbt run --select model_b+
- B: dbt run --select model_c+
- C: dbt run --select model_d
- D: dbt run --select model_d+

Consider this DAG showing 6 models for a production run. On the latest run, only model_d failed.

After fixing the underlying error, which command will build only the models that were not built on the recent run?

☐ dbt run --select model_a+

☐ dbt run --select model_b+

☐ dbt run --select model_d

☐ dbt run --select model_d+

3. Which is true about the doc function in dbt?
- A: The doc function can be used to reference the name of markdown file to be used in a description
- B: You can define a doc block in a SQL file and reference it in a description
- C: The doc function is used to create dependencies between models
- D: You can define a doc block within a YAML file and reference it in a description
- E: The doc function can be used to reference the name of a doc macro to be used in a description
- E: The doc function can be used to reference the name of a doc macro to be used in a description



4. Given the following models and their dependencies, the requirement is to check freshness before running any model. Choose 2 steps to achieve the requirement:

A: Create 2 jobs

B: Run source freshness --select source:source_1 and dbt build source:source_1+ sequentially in 1 job. Run dbt build source:source_2+ in another job

C: Run source freshness and dbt build sequentially

D: Create 1 job



C & D: Create 1 job, run source freshness and dbt build sequentially

5. Select "my_model" and its first-degree parents dbt run --select 1+my_model

6. When you've created a pull request and your teammate is reviewing your commits, is this situation in the development environment or the deployment environment?

Development Environment

7. Which feature of Jinja allows you to reuse code across multiple dbt models?

Macros

8. You are working on a large dbt project with multiple models. You have been asked to improve the readability of the database schema by using model aliasing. Which of the following statements about model aliasing is true?

B: Model aliasing can be used to override the default relation identifier for a model.



A: Model aliasing cannot be used in combination with custom schemas.

B: Model aliasing can be used to override the default relation identifier for a model.

C: Model aliasing is only useful for small projects with a few models.

D: Model aliasing cannot be used in a dbt project that uses package dependencies.

9. In dbt metrics, if you want to group your metrics by a custom dimension like "is_weekend", what variable must be set in the "dbt_project.yml" file?

A: is_weekend_dimension_list

B: custom_calendar_dimension_list

C: dimension_group_list

D: metric_group_list

B: custom_calendar_dimension_list

Explanation: To group your metrics by a custom dimension like "is_weekend" in dbt metrics, you need to set the variable "custom_calendar_dimension_list" in the "dbt_project.yml" file. This variable allows you to define custom calendar dimensions that can be used to group your metrics for reporting and analysis.

10. A data analyst wants to use snapshots in their dbt project to track changes in a dataset. What is the best practice for snapshotting the data in dbt?

A: Snapshot the data after it has been cleaned and transformed by downstream models.

B: Snapshot the data in its raw form and use downstream models to clean up the data.

C: Snapshot the data after it has been cleaned and transformed, but before it has

B: Snapshot the data in its raw form and use downstream models to clean up the data.



been joined with other tables.

D: C and A

-
11. What is a best practice in SQL for separating logic that cleans up data from logic that transforms data? A: Using common table expressions (CTEs)
- A: Using common table expressions (CTEs)
B: Using the ref function to build models on top of other models
C: Keeping all logic in a single query
D: Using the dbt tool to separate logic out into separate models
-
12. Alex is working on a dbt project and has defined several hooks for the models in both the `dbt_project.yml` file and in the config block of a specific model. He wants to know the order in which the hooks will be executed. How does dbt execute hooks defined in both the `dbt_project.yml` file and in the config block of a specific model? D: Both sets of hooks are executed, and the SQL statements defined in both will be executed in the order they were defined.
- A: Only the hooks defined in the config block of a specific model are executed.
B: Only the hooks defined in the `dbt_project.yml` file are executed.
C: Hooks are executed randomly, regardless of where they are defined.
D: Both sets of hooks are executed, and the SQL statements defined in both will be executed in the order they were defined.
-

13.



Alex is a data engineer working on a dbt project that uses the merge strategy with a defined `unique_key`. He notices that rows matching the `unique_key` are not being replaced with new data as expected. Which of the following reasons could explain this issue?

A: The `unique_key` is not set up correctly in the configuration.

B: The merge strategy is not suitable for handling unique keys.

C: The `unique_key` must be combined with another strategy for it to work correctly.

D: The `dbt_project.yml` file is missing or improperly configured.

A: The `unique_key` is not set up correctly in the configuration.

14. Do you need to create a target schema before running dbt?

A: Yes, you need to create the target schema before running dbt.

B: No, you do not need to create the target schema before running dbt. dbt will check if the schema exists when it runs, and create it if it does not exist.

C: You only need to create the target schema if you are running dbt on-premises, but not if you are using dbt Cloud.

D: It depends on the version of dbt you are using.

B: No, you do not need to create the target schema before running dbt. dbt will check if the schema exists when it runs, and create it if it does not exist.

15. You want to capture the current state of a source table in your dbt project at a specific

B: snapshot



point in time. Which dbt config should you use?

- A: seed
- B: snapshot
- C: archive
- D: model

-
16. You are an Analytics engineer working on a dbt project. Which of the following is a benefit of defining tests in dbt?
- A: Speed up the development process.
 - B: Allow for more flexibility in the code.
 - C: Improve the user interface of the project.
 - D: Confirm the code's functionality and prevent code regressions.
-
17. In a dbt project, how can you change the default location of the profiles.yml file to a different directory?
- A: Modify the dbt_project.yml file to include a reference to the new directory.
 - B: Use the DBT_PROFILES_DIR environment variable to change the default location of the profiles.yml file.
 - C: Include a --profiles-dir option in the dbt_project.yml file.
 - D: Move the profiles.yml file to the new directory and update the file path in each model's configuration.

D: Confirm the code's functionality and prevent code regressions.

B: Use the DBT_PROFILES_DIR environment variable to change the default location of the profiles.yml file.

Explanation: There are multiple ways to direct dbt to a different location for your profiles.yml file, but the most appropriate method is to use the DBT_PROFILES_DIR environment variable to change the default location.

Specifying this environment variable overrides the directory that dbt looks for your profiles.yml file in. You can specify this by running:

```
export DBT_PROFILES_DIR=path/to/directory.
```



18. **What is the reason why dbt Python models have limited access to the context of the project?**
- A: Python models do not support passing context as an argument to the `model()` function.
- B: Python models use a different templating language that does not have access to the project's context.
- C: Jinja, the templating language used by dbt SQL models, provides access to the project's context, which is not available to Python models.
- D: The context is not relevant or necessary for Python models to perform their functions.
- C: Jinja, the templating language used by dbt SQL models, provides access to the project's context, which is not available to Python models.
- Explanation: The reason why dbt Python models have limited access to the context of the project is that they do not use Jinja, a templating language that is used by dbt SQL models to render compiled code and provides access to the project's context. Instead, the context is made available from the dbt class and passed in as an argument to the `model()` function.
-
19. **In dbt, which method is used to configure a Python model within the model's .py file?**
- A: `dbt.configure()`
- B: `dbt.setup()`
- C: `dbt.config()`
- D: `dbt.define()`
- C: `dbt.config()`
-
20. **How can you specify a schema other than the target schema in your `profiles.yml` for building models in dbt, and where can you specify this information?**
- A: You can specify the schema in the `dbt_project.yml` file using the `model-paths` configuration.
- B: You can specify the schema in the model
- C: You can specify the schema in the `dbt_project.yml` file using the schema configuration block, or in the model file using a `config` block.



file using the ref function.

C: You can specify the schema in the dbt_project.yml file using the schema configuration block, or in the model file using a config block.

D: You can specify the schema in the model file using the schema keyword.

21. Which of the following is a valid use case for using variables in dbt models? D: Configuring time zones or avoiding hardcoded table names

A: Generating reports and visualizations from SQL queries

B: Storing configuration data for use in dbt projects

C: Defining data transformations in a graphical user interface

D: Configuring time zones or avoiding hardcoded table names

22. What will happen if you specify threads: 1 when running dbt? D: dbt will start building only one model, and finish it, before moving onto the next model, but the run time of the project will increase.

A: dbt will start building only one model, and finish it, before moving onto the next model.

B: dbt will start building all models at once, without finishing any of them.

C: dbt will start building all models in a sequential order

D: dbt will start building only one model, and finish it, before moving onto the next model, but the run time of the project will increase.

When you specify threads: 1 in your dbt project configuration, it means that you are restricting the number of threads that can be used to execute your dbt commands to 1. This implies that dbt will run only one operation at a time, i.e., build one model at a time, and wait for the completion of the current operation before proceeding to the next one. dbt will start building only one



model, and finish it, before moving onto the next model, but the run time of the project will increase.

23. Which of the following types of files are best suited for loading using dbt's seed functionality, according to dbt documentation?

A: Large CSV exports from a production database

B: Raw data files in JSON or XML format

C: Files containing business-specific logic, such as a list of country codes or employee IDs

D: Files containing sensitive information, such as PII or passwords

C: Files containing business-specific logic, such as a list of country codes or employee IDs

Explanation: According to dbt documentation, seeds are version controlled and maintainable, and are best suited to files containing business-specific logic, such as a list of country codes or employee IDs.

Seeds should not be used to load raw data, as this should be loaded using an ETL tool or SQL script, and large CSVs are not performant when loaded using dbt's seed functionality. Similarly, files containing sensitive information should not be stored in seeds, but instead should be secured and protected using appropriate security measures.

24. You are working on a dbt project where you have defined custom schemas for some models. During a dbt run, you notice that one of the models is not being created in the database. What is a possible reason for this issue?

A: The `generate_alias_name` macro has been overridden

B: There is an ambiguous model name in the schema

B & D: There is an ambiguous model name in the schema and/ or The target schema is not defined in the project configuration file

Explanation: One possible reason for a model not being created in the database is that there is an ambiguous model name in the schema. This can happen if the same model name is used for two or more models in different custom schemas. To resolve



C: The model does not have an alias defined

D: The target schema is not defined in the project configuration file

this issue, you can change the schema or alias configurations of the models to avoid ambiguity.

if the target schema is not defined in the project configuration file, dbt might not know where to create the model in the database. This could lead to the model not being created as expected.

25. You are a data scientist working on a data pipeline and you want to write a model in dbt that doesn't directly build into the database. Which materialization option should you use for this purpose?

A: Table

B: View

C: Incremental

D: Ephemeral

D: Ephemeral

Explanation: The ephemeral materialization in dbt allows you to write reusable logic without directly building the model into the database. Instead, the code from this model will be used as a common table expression in dependent models. This materialization option is best suited for very lightweight transformations that don't need to be directly queried, and are used only in one or two downstream models. If you want to write a model in dbt that doesn't directly build into the database, you should use the ephemeral materialization.

26. As a data engineer, you are concerned about securely storing credentials for your dbt project. What method is considered more secure for storing credentials in dbt?

A: Storing credentials directly in the profiles.yml file.

B: Using environment variables to load cre-

B: Using environment variables to load credentials.



dentials.

C: Saving credentials in a separate plain-text file within the project directory.

D: Storing credentials in the dbt_project.yml file.

27. What is the purpose of the `--vars` command line option in dbt? B: To override variables for a run of dbt

A: To define variables in a dbt project

B: To override variables for a run of dbt

C: To configure timezones and avoid hard-coded table names

D: To generate reports and visualizations from SQL queries

28. When implementing data integrity checks in dbt, which of the following methods should be avoided? B: Running data integrity tests against individual models

A: Running data integrity tests against the data warehouse

B: Running data integrity tests against individual models

C: Automating data integrity checks

D: Performing manual checks on the data

E: None of the above

29. Which of the following statements are true about the `ref` function in dbt? D: All of the above

A: The `ref` function allows you to build models on top of other models

B: The `ref` function allows you to modularize your models



C: The ref function allows you to make your models reusable
D: All of the above

30. A company wants to use the check strategy to detect changes in their product inventory. Which configuration must be specified to use the check strategy in dbt? B: check_cols
- A: updated_at
B: check_cols
C: unique_key
D: None of the above
-

31. Which of the following best describes how dbt optimizes your workflow? C: By building up reusable, modular data models
- A: By creating a new database with each dbt run
B: By processing raw data in real-time
C: By building up reusable, modular data models
D: By storing data in a more efficient format like avro and parquet
-

32. A team is working on a dbt project where they are using snapshots to track changes to their product inventory. They want to prevent accidental deletion of their snapshots. What is the recommended approach for setting privileges on the snapshot tables? B: Run the snapshots as a different user or role in the warehouse
- A: Run the snapshots as the same user or role as other models in the warehouse



B: Run the snapshots as a different user or role in the warehouse

C: Don't set any privileges on the snapshots

D: Give all users full access to the snapshots

33. When using dbt to test a model, what does dbt do with the queries it constructs for each test?

C: It executes them and returns the number of records that failed the test.

A: It saves them to the YAML file for future reference.

B: It discards them after the test is run.

C: It executes them and returns the number of records that failed the test.

D: It compiles them into a single executable query.

34. What command should you use to initiate a new dbt project?

B: dbt init dbt_project_name

A: dbt new dbt_project_name

B: dbt init dbt_project_name

C: dbt create dbt_project_name

D: dbt start dbt_project_name

35. Sarah is a data analyst who works with a dbt Cloud project. She needs to define an environment variable to store the database password. Which prefix should she use for the variable name to define it as an environment secret variable in dbt Cloud?

C: DBT_ENV_SECRET

A: DBT



B: DBT_ENV

C: DBT_ENV_SECRET

D: DBT_SECRET

36. You are reviewing a dbt project that has a source schema defined in the schema.yml file. The source schema has a name in the database that looks like a machine generated name which is not easily readable, you want to use a more sensible name in dbt. How can you implement this and still ensure dbt compiles using the actual source schema name?
- A: Use DDL and DML statements to change the schema name to one that can be easily read and makes sense
- B: Use the schema property to define the actual names as per the database, and use your name: property for the name that makes sense and is easily readable
- C: Use the 'database' property to define the actual names as per the database, and use your name: property for the name that makes sense and is easily readable
- D: Use the 'table' property to define the actual names as per the database, and use your name: property for the name that makes sense and is easily readable
37. As a data engineer working on a dbt project, you need to install a package from a tarball URL hosted internally within your
- B: Use the schema property to define the actual names as per the database, and use your name: property for the name that makes sense and is easily readable
- B: The tarball URL and the subfolder name where the package source code is installed.



organization. What information do you need to provide in order to install this package?

A: The tarball URL and the package version.

B: The tarball URL and the subfolder name where the package source code is installed.

C: The tarball URL and the branch name of the package repository.

D: The tarball URL and the name of the data warehouse used in the project.

Explanation: To install packages using tarball URLs hosted internally, you need to specify the URL of the tarball and the subfolder name where the package source code is installed within. This information will allow dbt to fetch and install the package from the internal location, ensuring that the correct version and source code are used in your project.

38. Maria is working on a dbt project and needs to create snapshots for a table that does not have a reliable `updated_at` column. Which snapshot strategy should she use to compare a list of columns between their current and historical values?

A: Timestamp strategy

B: Merge strategy

C: Check strategy

D: Incremental strategy

C: Check strategy

Explanation: Use the check strategy for tables that do not have a reliable `updated_at` column. This strategy works by comparing a list of columns between their current and historical values. If any of these columns have changed, then dbt will invalidate the old record and record the new one. If the column values are identical, then dbt will not take any action.

39. You are running dbt and want to check the connection to your warehouse. What should you do?

A: Use the debug command within the dbt project

B: Use the test command within the dbt project

C: Use the run command within the dbt

A: Use the debug command within the dbt project

Simply run `dbt debug` from within dbt project to test your connection.



project

D: Use the check command within the dbt project

40. What is the purpose of the dbt analysis folder?

A: To store SQL queries for use in dbt models.

B: To store SQL queries for use in dbt tests.

C: To organize analytical SQL queries in a dbt project.

D: To store SQL queries for use in dbt seeds.

C: To organize analytical SQL queries in a dbt project.

Explanation: These queries are not run by dbt when you initiate the dbt run command, but can be compiled using the dbt compile command. This feature provides a way to version control your non-model SQL queries and keep them organized alongside your other dbt project components. Option A is incorrect because SQL queries for use in dbt models are stored in the models folder. Option B is incorrect because SQL queries for use in dbt tests are stored in the tests folder. Option D is incorrect because SQL queries for use in dbt seeds are stored in the seeds folder.

41. A data analyst wants to create a snapshot configuration to track changes in a customer table that contains sensitive information. They want to ensure that the snapshot configuration provides the best handling of changes to the source data, including the addition and deletion of columns. Which strategy should they use when working with snapshot tables in dbt?

A: The timestamp strategy

B: The check_cols strategy

A: The timestamp strategy

Explanation: The timestamp strategy provides better handling of changes to the source data, including additions and deletions of columns, compared to the check_cols strategy. Using the timestamp strategy is considered a best practice for snapshot configurations.



C: The unique key strategy
D: The hard delete strategy

-
42. Which of the following statements is true about the Python model in dbt? B: The Python model is stored in a .py file in the models/ folder.
- A: The Python model is stored in a database table.
- B: The Python model is stored in a .py file in the models/ folder.
- C: The Python model is a class compiled by dbt Core.
- D: The Python model defines a function named dbt().
-
43. How can using dbt sources improve data analysis workflows? A & B: It allows for easy dependency tracing, It allows for source freshness reporting
- Select all answers that apply:
- A: It allows for easy dependency tracing
- B: It allows for source freshness reporting
- C: It allows for better data visualization
- D: It allows for easier data cleansing
-
44. A data team wants to reduce the number of columns that need to be checked in their database when using the check strategy in dbt. What is the recommended technique they can use to achieve this? C: Use a surrogate key to condense many columns into a single column
- A: Include all columns in the check_cols configuration
- B: Exclude columns from the check_cols configuration
- C: Use a surrogate key to condense many



columns into a single column

D: None of the above

-
45. You have created a metric that uses multiple dimensions, but when you run the metric, the resulting number is not providing useful information. What is the most likely reason for this issue?
- A: The metric is not defined correctly in the dbt project.
- B: The dimensions being used in the metric calculation are not relevant to the data being analyzed.
- C: The table being used to calculate the metric does not contain enough data to make the calculation meaningful.
- D: The metric is being calculated using an incorrect formula.
- B: The dimensions being used in the metric calculation are not relevant to the data being analyzed.
- Explanation: When a metric is created in dbt that uses multiple dimensions, it takes all the information in the table and simplifies it into a single number that gives you a quick understanding of the data.
-
46. What is the command used to specify a target other than the default when using dbt?
- A: --target
- B: --prod
- C: --dev
- D: --target_schema
- A: --target
-
47. In a dbt Python model, what is the concept of "lazy evaluation"?
- A: The ability to preview data using methods like `.show()` or `.head()` in development.
- B: The ability to execute Python code re-
- D: The output of each dataframe operation not being immediately calculated, but only computed when explicitly asked for.
- Explanation: In a dbt Python model, the concept of "lazy evaluation" refers to the



motely on a data platform.

C: The ability to create a series of meaningful transformations using CTEs.

D: The output of each dataframe operation not being immediately calculated, but only computed when explicitly asked for.

output of each dataframe operation not being immediately calculated. Instead, the operations are only computed when you explicitly ask for the final result of the data. In development, you can preview the data using methods like `.show()` or `.head()`. When you run a Python model, the full result of the final DataFrame will be saved as a table in your data warehouse.

48. **How can you specify column types in dbt?**

A: Use the type configuration block in the `dbt_project.yml` file to specify the column types for each model.

B: Use the `dtype` function to specify column types in the SQL select statement.

C: Cast the column to the correct type in your model using SQL syntax.

D: Use the `column_type` configuration block in the model file to specify the column types for each column.

C: Cast the column to the correct type in your model using SQL syntax.

49. **What happens when dbt encounters an error in the SQL code of a model during execution?**

A: dbt returns the error generated by the database or data warehouse, and downstream models are skipped.

B: dbt continues running the remaining models and ignores the model with the error.

C: dbt automatically fixes the error and re-

A: dbt returns the error generated by the database or data warehouse, and downstream models are skipped.



runs the model.

D: dbt stops execution immediately and does not provide any error message.

50. You are a data platform engineer working with a Python model that requires incremental updates. What materialization options do you have for Python models in dbt?

A: View and incremental

B: Table and view

C: Table and incremental

D: Ephemeral and table

C: Table and incremental

Explanation: In dbt, Python models have two materialization options: table and incremental. Incremental models in Python support the same incremental strategies as SQL models, but the specific strategies depend on the database adapter being used. It is not possible to use view or ephemeral materialization for Python models, nor can Python be used for non-model resources like tests and snapshots. Therefore, for a Python model that requires incremental updates, you can use either the table or incremental materialization. The incoming data must be filtered to only include new rows. The insert_overwrite strategy for incremental models is not yet supported for BigQuery/Dataproc, but the merge incremental strategy is supported.

51. Imagine you are a data engineer working for a company that deals with massive amounts of data. You have been tasked with optimizing the data transformation process to reduce the time it takes to process the data. You have been using dbt as your primary tool for data transformation, but you have noticed that the process

B: The "incremental_predicates" feature in dbt optimizes the incremental build process for large data volumes by using specified SQL expressions to determine changed data and update only those records, resulting in improved performance and reduced processing time for data transformations.



is taking too long when working with large data volumes. You have heard about a feature in dbt called "incremental_predicates" that can help improve performance in such scenarios. How does this dbt feature does this?

A: The "incremental_predicates" feature in dbt optimizes the incremental build process for large data volumes by using Tree based algorithms to determine changed data and update only those records, resulting in improved performance and reduced processing time for data transformations.

B: The "incremental_predicates" feature in dbt optimizes the incremental build process for large data volumes by using specified SQL expressions to determine changed data and update only those records, resulting in improved performance and reduced processing time for data transformations.

C: The "incremental_predicates" feature in dbt optimizes the incremental build process for large data volumes by using specified Jinja control structures to determine changed data and update only those records, resulting in improved performance and reduced processing time for data transformations.

D: dbt has no incremental_predicates feature, dbt uses pre-hook and post-hooks to

Explanation: The "incremental_predicates" feature in dbt is used to improve performance when working with large data volumes. It allows for a list of SQL expressions to be specified that will be used to optimize the incremental build process. By specifying these expressions, dbt can better determine which data has changed and only update those records instead of reprocessing the entire dataset. This can significantly improve performance and reduce the time it takes to run data transformations on large datasets.



determine changed data and update only those records, resulting in improved performance and reduced processing time for data transformations.

52. You have an incremental model in dbt with the default setting "on_schema_change: ignore", but when running the dbt command, the new column you added is not appearing in the target table. What could be the cause of this issue?
- A: The database does not support adding new columns in this way.
- B: The on_schema_change parameter is not set correctly in the configuration.
- C: A full-refresh of both the incremental model and any related models has not been executed.
- D: There is an issue with the SQL used in the model.
- C: A full-refresh of both the incremental model and any related models has not been executed.
- Explanation: When using an incremental model in dbt with the "on_schema_change: ignore" setting, any changes made to the columns in the incremental model will not be reflected in the target table until a full-refresh of both the incremental model and any related models is executed. This means that if you add a column to the incremental model, it will not appear in the target table until a full-refresh is executed. Therefore, the most likely cause of this issue is that a full-refresh of both the incremental model and any related models has not been executed.
53. How do exposures help data consumers understand and use the outputs of a dbt project?
- A: By providing a dedicated page in the auto-generated documentation site with context relevant to the outputs.
- B: By creating a separate schema for each user to maintain separate development and production environments.
- A: By providing a dedicated page in the auto-generated documentation site with context relevant to the outputs.
- Explanation: Exposures in dbt help data consumers understand and use the outputs of a project by providing a dedicated page in the auto-generated documentation site with context relevant to the outputs. This allows data consumers to eas-



C: By enabling you to test and run resources that feed into the exposure.

D: By defining and describing the upstream use of the project.

ily access information about the outputs and how they can be used in downstream applications, dashboards, or data science pipelines.

54. Select the best reason for using dbt sources rather than referencing raw database tables choose all answers that apply:
- A: It allows for easy dependency tracing
 - B: It improves performance of SQL queries
 - C: It allows for version control of data sources
 - D: It improves data visualization capabilities

A & C: It allows for easy dependency tracing, It allows for version control of data sources

55. What are the limitations of the `dbt.config()` method in terms of the types of arguments that can be passed to it?
- A: It can accept any type of data structure or function.
 - B: It can only accept strings and numbers.
 - C: It can only accept basic data types such as strings, booleans, and numbers.
 - D: It can accept any data type as long as it is defined in a YAML file.

C: It can only accept basic data types such as strings, booleans, and numbers.

56. A healthcare company wants to keep a historical record of changes to their patient data, such as changes to their medical history and medication information. which feature of dbt is the best for this scenario?
- A:** dbt historical cannot be used for tracking

B: dbt snapshots can be used to keep a historical record of changes to patient data, including medical history and medication information.

By using snapshots, the healthcare company can implement type-2 Slowly Chang-



changes in patient data as it violates HIPAA regulations.

B: dbt snapshots can be used to keep a historical record of changes to patient data, including medical history and medication information.

C: dbt data lineage can only be used for tracking changes to dimension tables in data warehousing and cannot be used in healthcare settings.

D: dbt incremental can be used for tracking changes to patient data but only if the patient has given their explicit consent.

ing Dimensions (SCDs) to track changes to rows in mutable source tables over time, making it easier to track the evolution of patient data over time.


57. You are working on a dbt project and want to select models that have a tag of "marketing" and are part of the "core" package. Which of the following YAML selectors will achieve this?



A Selector A uses the 'intersection' clause to select models that meet both conditions: being part of the 'core' package and having a 'marketing' tag. This fulfills the requirement of selecting models with a tag of "marketing" and part of the "core" package.

58. You want to select models in the "sales" package that have a materialization setting of "incremental" but exclude those tagged as "deprecated". Which of the following YAML selectors will achieve this?

B Selector B uses the 'intersection' clause to select models that are part of the 'sales' package and have a materialization setting of 'incremental'. Then, it uses the 'exclude' clause to remove models with a 'deprecate-



ed' tag from the selection. This meets the requirements of the question.

59. The marketing team later requests that you exclude a specific model, customer_lifetime_value, while still executing the other customer_analytics tagged models and their dependencies. How would you modify your dbt command to accommodate this request?

B: dbt run --select tag:customer_analytics+
--exclude customer_lifetime_value

- A: dbt run --select tag:customer_analytics+
--select -customer_lifetime_value
B: dbt run --select tag:customer_analytics+
--exclude customer_lifetime_value
C: dbt run --select tag:customer_analytics+
--select -+customer_lifetime_value
D: dbt run --select tag:customer_analytics+
-customer_lifetime_value

60. You have been asked to run a specific set of models related to the daily financial reports, but you need to exclude models tagged as "nightly" and those with a materialized configuration set to "table." What command would you use to run the models contained in the "daily_reports" directory, excluding those with the "nightly" tag and materialized as "table" ?

A: dbt run --select path:models/daily_reports --exclude tag:nightly --exclude config.materialized:table

- A: dbt run --select path:models/daily_re-



ports --exclude tag:nightly --exclude config.materialized:table

B: dbt run --select path:models/daily_reports --select -tag:nightly,config.materialized:table

C: dbt run --select path:models/daily_reports,tag:nightly,config.materialized:table

D: dbt run --select path:models/daily_reports --select -tag:nightly --select -config.materialized:table

61. What is the main purpose of the `source_status` method in dbt?
- A: To calculate the total execution time of dbt commands
- B: To evaluate the freshness of data sources and apply selectors based on the result
- C: To optimize query performance
- D: To manage user permissions and access control

B: To evaluate the freshness of data sources and apply selectors based on the result

Explanation: The `source_status` method is used to evaluate the freshness of data sources by comparing the current state with the previous state. It allows users to apply selectors based on the freshness of the sources, which can be useful in ensuring that the data used in the transformation is up-to-date.

62. Which of the following dbt commands produces a `sources.json` artifact that can be referenced in subsequent dbt invocations?
- A: dbt source new_source
- B: dbt test freshness
- C: dbt source freshness
- D: dbt run freshness

C: dbt source freshness

Explanation: The dbt source freshness command produces a `sources.json` artifact that contains information about the execution times and `max_loaded_at` dates for dbt sources. This artifact can be used to reference the source freshness results in subsequent dbt invocations.

63.



What is the primary function of the dbt source freshness command?

A: To generate a list of all available data sources

B: To update the data sources to their latest versions

C: To evaluate the freshness of the data sources

D: To create a new source for the dbt project

C: To evaluate the freshness of the data sources

64. How can you reference the source freshness results in a subsequent dbt command?

A: Use the --source-status flag

B: Use the --select flag with the source_status selector

C: Use the --refresh flag

D: Use the --state flag with the DBT_ARTIFACT_STATE_PATH environment variable

B: Use the --select flag with the source_status selector

Explanation: To reference the source freshness results in a subsequent dbt command, you can use the --select flag followed by the source_status selector (e.g., source_status:fresher+). This allows you to apply selectors based on the freshness of the data sources.

65. You are working on a dbt project and want to ensure that only the data sources with freshness greater than the previous state are used for the transformation. Which selector should you use in your dbt command?

A: source_status:fresher

B: source_status:fresher+

C: source_status:stale

D: source_status:stale+

B: source_status:fresher+

Explanation: To select only the data sources with freshness greater than the previous state, you should use the 'source_status:fresher+' selector in your dbt command. This selector will ensure that the transformation only uses data sources that are fresher than their previous state.



-
66. In a dbt project, you need to build a model using data sources that are no older than their previous state. Which selector should you use to achieve this?
- A: `source_status:fresher`
B: `source_status:fresher+`
C: `source_status:equal`
D: `source_status:stale`
- A: `source_status:fresher`
Explanation: To build a model using data sources that are no older than their previous state, you should use the 'source_status:fresher' selector in your dbt command. This selector will ensure that the transformation only uses data sources that are either fresher or equal to their previous state.
-
67. A data analyst in your team wants to analyze data sources that have become stale compared to their previous state. Which selector should you use in your dbt command to meet this requirement?
- A: `source_status:fresher`
B: `source_status:fresher+`
C: `source_status:stale`
D: `source_status:stale+`
- C: `source_status:stale`
Explanation: To analyze data sources that have become stale compared to their previous state, you should use the 'source_status:stale' selector in your dbt command. This selector will enable the data analyst to focus on the data sources that have become less fresh compared to their previous state.
-
68. You are working on a dbt project and need to exclude data sources that have become stale compared to their previous state. Which selector should you use in your dbt command?
- A: `source_status:exclude_stale`
B: `source_status:fresher`
C: `source_status:stale`
D: `source_status:fresher+`
- B: `source_status:fresher`
Explanation: To exclude data sources that have become stale compared to their previous state, you should use the 'source_status:fresher' selector in your dbt command. This selector will ensure that the transformation only uses data sources that are either fresher or equal to their previous state, effectively excluding the stale ones.
-
69. In a dbt project, you want to select data sources that are as fresh as their previous state, but not fresher. Which selector
- A: `source_status:equal`
Explanation: To select data sources that are at least as fresh as their previous



should you use in your dbt command?

- A: `source_status:equal`
- B: `source_status:fresher`
- C: `source_status:stale`
- D: `source_status:fresher+`

state but not fresher, you should use the 'source_status:equal' selector in your dbt command. This selector will ensure that the transformation only uses data sources that have the same freshness as their previous state.

70. You have this DAG showing the 3 models, how do you ensure you run just the stg_covid_cases model?

- A: `dbt run --s stg_covid_cases`
- B: `dbt run --select model=+stg_covid_cases+`
- C: `dbt run -s stg_covid_cases`
- D: `dbt run -select stg_covid_cases`



C: `dbt run -s stg_covid_cases`

71. You initiated a dbt run as seen in the diagram, How do you fix the cycle error encountered?

- A: Run the dbt self-heal command.
- B: Rerun the model using `dbt run --full-refresh`
- C: Update the ref functions to remove the cycle.
- D: Rerun the model using `dbt run --vars '{"cyclic": "false"}'`



C: Update the ref functions to remove the cycle.

Explanation: We observe cyclic errors in dbt when there is a circular dependency between two or more models, causing an infinite loop during the build process. This happens when a model depends on another model, which in turn depends on the first model, creating a cycle in the dependency graph. It can also occur with indirect dependencies involving more than two models. To fix this we need to update the ref function which builds dependencies between model and break the cycle.



-
72. Which of the following statements are true about developing in dbt? Options:
- A: Data models are saved as tf(transformation) files.
- B: Data models are saved as sql files.
- C: Data models are saved as yaml files.
- D: Data models are saved as json files.
- B: Data models are saved as sql files.
-
73. Where can variables be defined in a dbt project?
- A: In the data directory
- B: In the models directory
- C: In the macros directory
- D: In the dbt_project.yml file or on the command line
- D: In the dbt_project.yml file or on the command line
- Explanation: Variables can be defined in the dbt_project.yml file or on the command line. When defined in the dbt_project.yml file, variables can be set globally for the entire project, and can be referenced in any model, analysis, test, or hook. When defined on the command line, variables can be set on a per-run basis and can override variables defined in the dbt_project.yml file. This can be useful for setting variables that are specific to a particular run or environment.
-
74. You have defined an exposure called "my_exposure" in your dbt project, but when you try to run the dbt run command with the -s flag and the exposure name, you receive an error message that says "Could not find exposure 'my_exposure'". What is the most likely reason for this error?
- A: The exposure has not been defined correctly in the dbt schema.yml file.
- A: The exposure has not been defined correctly in the dbt schema.yml file.



rectly in the dbt schema.yml file.

B: The exposure has not been added to the dbt_project.yml file.

C: The exposure has not been added to the auto-generated documentation site.

D: The exposure has not been tested and run in the dbt project.

75. dbt Exposures

Exposures make it possible to define and describe a downstream use of your dbt project, such as in a dashboard, application, or data science pipeline. By defining exposures, you can then:

- run, test, and list resources that feed into your exposure
- populate a dedicated page in the auto-generated documentation site with context relevant to data consumers

76. When working with dbt, you want to run a specific model. How can you accomplish this?

A: Use the --select flag followed by the name of the model

B: Use the -s flag followed by the name of the model

C: Use the --run flag followed by the name of the model

D: Use the -r flag followed by the name of the model

A&B: Use the --select flag followed by the name of the model, Use the -s flag followed by the name of the model

dbt run --select Model_Name

dbt run -s Model_Name

77.

C: vars



You want to define a custom project variable in the `dbt_project.yml` file that can be used for data compilation. What configuration option should you use?

- A: macro-paths
 - B: test-paths
 - C: vars
 - D: analysis-paths
-

78. What is the order of precedence for variable declaration in dbt?
- A: Variables defined with `--vars`, global variable declaration in the `dbt_project.yml` file, package-scoped variable declaration in `dbt_project.yml` file, and the variable's default argument.
- B: Package-scoped variable declaration in `dbt_project.yml` file, variables defined with `--vars`, global variable declaration in the `dbt_project.yml` file, and the variable's default argument.
- C: Variables defined with `--vars`, package-scoped variable declaration in `dbt_project.yml` file, global variable declaration in the `dbt_project.yml` file, and the variable's default argument.
- D: Global variable declaration in the `dbt_project.yml` file, variables defined with `--vars`, package-scoped variable declaration in `dbt_project.yml` file, and the variable's default argument.
- C: Variables defined with `--vars`, package-scoped variable declaration in `dbt_project.yml` file, global variable declaration in the `dbt_project.yml` file, and the variable's default argument.
- Explanation: In dbt, the order of precedence for variable declaration is: variables defined with `--vars` command line argument, package-scoped variable declaration in `dbt_project.yml` file, global variable declaration in the `dbt_project.yml` file, and the variable's default argument. If dbt is unable to find a definition for a variable, a compilation error is raised.
-



79. Which method is used to set the configurations within a Python model in dbt? B: dbt.config()
- A: dbt.setup()
 - B: dbt.config()
 - C: dbt.define()
 - D: dbt.run()
-
80. Jack is working on a dbt project and wants to change the default identifier for a relation that is created when a model is executed. What model configuration option should he use to achieve this? C: The 'alias' configuration
- A: The 'name' configuration
 - B: The 'identifier' configuration
 - C: The 'alias' configuration
 - D: The 'relation_name' configuration
-
81. A data engineer is working on an incremental model in dbt using the delete+insert+merge strategy. However, the incremental model run is failing with a unique key violation error. What could be the reason for this error? C: The unique_key is not truly unique in the data.
- A: The unique_key is not defined in the model.
 - B: The unique_key is not used in the delete+insert+merge strategy.
 - C: The unique_key is not truly unique in the data.
 - D: The delete+insert+merge strategy is not appropriate for the data.
-



82. Which of the following statements about dbt are true? A, B, C, D, E, F: All are true

Choose all that apply:

A: dbt helps you get more work done

B: dbt produces higher quality results

C: dbt allows you to modularize and centralize your analytics code

D: dbt provides guardrails typically found in software engineering workflows

E: dbt allows you to collaborate, version, test, and document your queries before deploying them to production

F: dbt provides monitoring and visibility

83. You are a data analyst working at a company that utilizes dbt cloud. You need to schedule a job to run daily at 8 AM that will compile and test your company's data models. Which method should you use to accomplish this task? B: Use the dbt cloud console to create and run the job daily at 8 AM

A: Use the dbt run command in dbt core and set the schedule according to your organization's best practices

B: Use the dbt cloud console to create and run the job daily at 8 AM

C: Use the dbt test command in dbt core and set the schedule according to your organization's best practices

D: Use the dbt compile command in dbt core and set the schedule according to your organization's best practices



-
84. You have created a metric in your dbt project that uses one dimension - time. However, when you run the metric, the resulting number does not match the expected value. What is the most likely reason for this issue?
- A: The metric is not defined correctly in the dbt project.
- B: The time dimension is not being properly considered in the metric calculation.
- C: The table being used to calculate the metric does not contain accurate data.
- D: The metric is using too many dimensions to accurately represent the data.
- C: The table being used to calculate the metric does not contain accurate data.
- Explanation: When a metric is created in dbt, it takes all the information in a table and simplifies it into a single number that gives you a quick understanding of the data. If the resulting number does not match the expected value, the most likely reason is that the table being used to calculate the metric does not contain accurate data. This can be due to incorrect data being loaded into the table, data being filtered incorrectly, or other issues with the underlying data.
-
85. Sarah is working on a dbt project that includes Python models. Which materialization options are available for these Python models?
- A: View and Ephemeral
- B: Table and Ephemeral
- C: Table and Incremental
- D: View and Incremental
- C: Table and Incremental
- Explanation: Python models in dbt cannot be materialized as view or ephemeral. They can only be materialized as table (the default materialization) or incremental.
-
86. Anna is considering using dbt for her data transformation project and wants to understand the differences between dbt Core and dbt Cloud. What are some of the benefits of using dbt Cloud over dbt Core?
- A: Lower cost and more extensive customization options
- B: Web-based UI, hosted environment, and differentiated features



differentiated features

C: Access to a larger open-source community and more extensive documentation

D: Higher performance and greater compatibility with various data warehouses

87. You are a dbt developer working on a lightweight transformation that is only used in one downstream model and don't need to be directly queried. Which materialization should you use for this model?

A: View

B: Table

C: Incremental

D: Ephemeral

D: Ephemeral

The ephemeral materialization in dbt allows you to write reusable logic without directly building the model into the database. Instead, the code from this model will be used as a common table expression in dependent models. This type of materialization has advantages such as keeping the data warehouse clean by reducing clutter, and allowing for lighter weight transformations. Therefore, for a lightweight transformation that is only used in one downstream model and doesn't need to be directly queried, you should use the ephemeral materialization. However, it also has some limitations, such as not being able to select directly from the model and some operations not being able to reference it.

88. Sara is a data analyst working with a dbt Cloud project. She has noticed that her dbt project is not compiling to the values she has set. What can she do to resolve this issue?

A: Change the variable values to their de-

D: Delete the target/partial_parse.msg-pack file in the project.

Explanation: In dbt Cloud, partial parsing of a project can lead to issues. If you find that your dbt project is not compiling to the values you have set, deleting the



fault settings.

B: Rename the environment variables in the dbt Cloud UI.

C: Delete the dbt Cloud project and create a new one.

D: Delete the target/partial_parse.msgpack file in the project.

target/partial_parse.msgpack file in your project can help. Doing so will force dbt to recompile your entire project and may help resolve any issues caused by partial parsing.

89. A data team is setting up a new dbt project to transform their company's employee data. They have a CSV file with a list of department codes and descriptions that they want to load into their database as reference data. Which of the following is the correct approach for this scenario?

A: Use a dbt snapshot to load the department data into the database

B: Write a custom Python script to load the department data into the database

C: Use a dbt seed to load the department data into the database

D: Write a set of SQL queries to insert the department data into the database

C: Use a dbt seed to load the department data into the database

Explanation: To load a CSV file with reference data into a database in a dbt project, the correct approach is to use a dbt seed. This allows the data team to load the department data into the database and use it in their transformation logic.

90. You want to add tests to a model in a project. You want to ensure that your models are working correctly by testing individual columns in the models. What should you do?

A: Create a Schema YAML file for that specific model in the models directory

B: Add the test type you want into the YAML

A, B & C: Create a Schema YAML file for that specific model in the models directory, Add the test type you want into the YAML file, Run the dbt test command, and confirm that all your tests passed



file.

C: Run the dbt test command, and confirm that all your tests passed.

D: Run the dbt test command, and confirm that all your tests failed.

-
91. In dbt, what is the recommended practice when developing locally?
- A: Use the "dev" target as the default
B: Use the "prod" target as the default
C: Use the "--target" option when issuing a DBT command
D: Use the "default" target when issuing a DBT command
- A: Use the "dev" target as the default
Explanation: Dbt Supports Multiple Targets Within A Profile, Which Encourages The Use Of Separate Development And Production Environments. When Developing In Dbt Locally It Is Always Good Practice To Use The Dev Target Usually Set As The Default. A Separate "Prod" Target Can Also Be Created For Production Environments. Users Can Also Use The "--Target" Option When Issuing A Dbt Command To Use A Target Other Than The Default
-
92. In dbt, what are the three ways to configure dbt Python models?
- A: In a SQL file, in a .py file, and in the dbt_project.yml file.
B: In a dedicated .yml file, within the models/ directory, and in a SQL file.
C: In dbt_project.yml, in a dedicated .yml file within the models/ directory, and within the model's .py file using the dbt.config() method.
D: In a dedicated .py file, within the models/ directory, and within the model's .yml file using the dbt.config() method.
- C: In dbt_project.yml, in a dedicated .yml file within the models/ directory, and within the model's .py file using the dbt.config() method.
Explanation: The three ways to configure dbt Python models in dbt are:
1) in dbt_project.yml, where you can configure many models at once,
2) in a dedicated .yml file, within the models/ directory, and
3) within the model's .py file, using the dbt.config() method. Therefore, option C is



correct, while the other options are incorrect.

93. One of your team members is new to using dbt and is unsure how to reference a source in a model. How would you explain this process to them?

A: Once a source has been defined in the schema.yml file, you can reference it in a model using the `{{ source() }}` function. This function allows you to select data from the source in the same way you would select data from a table in an SQL query.

B: Once a source has been defined in the schema.yml file, you can reference it in a model using the `{{ model() }}` function. This function allows you to select data from the source in the same way you would select data from a table in an SQL query.

C: Once a source has been defined in the schema.yml file, you can reference it in a model using the `{{ db() }}` function. This function allows you to select data from the source in the same way you would select data from a table in an SQL query.

D: Once a source has been defined in the schema.yml file, you can reference it in a model using the `{{ join() }}` function. This function allows you to combine data from different sources in the same way you would combine data from different tables in an SQL query.

A: Once a source has been defined in the schema.yml file, you can reference it in a model using the `{{ source() }}` function. This function allows you to select data from the source in the same way you would select data from a table in an SQL query.



-
94. You are working on a dbt python project that involves data processing and transformation. You have specified specific package versions in the dbt_project.yml file under the "packages" configuration block. However, when you run dbt, you notice that some models are still using the wrong package versions. Which of the following could be the reason for this?
- A: The specified package versions are not compatible with the data platform being used.
- B: The specified package versions are not installed on the system.
- C: The package versions are not being read from the dbt_project.yml file.
- D: The models are not referencing the correct package names.
- B: The specified package versions are not installed on the system.
- Explanation: If some models are still using the wrong package versions, it is possible that the specified package versions are not installed on the system. It is important to ensure that the specified packages and their versions are installed and accessible to the Python runtime on your data platform.
-
95. In dbt, which of the following methods returns a DataFrame pointing to a upstream model?
- A: dbt.run()
- B: dbt.source()
- C: dbt.ref()
- D: dbt.snapshot()
- C: dbt.ref()
-
96. You are a dbt dev working on a dbt project. You want to make a one-off assertion about a specific resource in your project. Which type of test should you create?
- C: Singular test
- Explanation: Singular tests in dbt are a way to make one-off assertions about specific resources in your dbt project such as mod-



- A: Generic test
- B: Built-in test
- C: Singular test
- D: Parameterized test

els, sources, seeds, and snapshots. These tests are defined in .sql files located in the tests directory specified in the test-paths configuration. To make a one-off assertion about a specific resource in your project, you should create a singular test. Singular tests are easy to create and can help ensure the integrity of your dbt project.

97. What is the minimum requirement for a dbt project?
- A: A dbt_project.yml project configuration file
 - B: A models directory with transformation logic
 - C: A connection to a database
 - D: A set of predefined macros

A: A dbt_project.yml project configuration file
Explanation: The minimum requirement for a dbt project is the presence of a dbt_project.yml project configuration file. This file provides the necessary configuration for dbt to run commands and connect to the appropriate database. Options B, C, and D are not required for a dbt project to function.

98. You are a data engineer working on a dbt project. You are in charge of creating models that will be used by different teams in your organization. You have created a model named "customer_orders.sql" and you want to provide a more descriptive name for the relation in the database. Which of the following is a valid way to do this?
- A: Change the filename of the model to a more descriptive name
 - B: Add an alias configuration to the model's

B: Add an alias configuration to the model's configuration with a custom value
Explanation: To provide a more descriptive name for the relation in the database, you can use the alias model configuration. This allows you to override the default relation name, which is based on the model's filename.



configuration with a custom value

C: Use the `generate_schema_name` macro to set a custom name for the relation

D: Use the `rename_relation` function in the model code to change the name of the relation

99. In order for the migration of legacy SQL code to dbt to work without any problems, which of the following needs to be taken into consideration?
- D: All of the above
- A: The legacy SQL dialect
- B: Stored procedures or functions
- C: The data warehouse
- D: All of the above

100. By default, where does dbt expect model files to be located in a dbt project, and how can you change this?
- B: By default, dbt expects model files to be located in the `models` subdirectory of your project. You can change this by updating the `model-paths` configuration in your `dbt_project.yml` file.
- A: By default, dbt expects model files to be located in the project root directory. You cannot change this.
- B: By default, dbt expects model files to be located in the `models` subdirectory of your project. You can change this by updating the `model-paths` configuration in your `dbt_project.yml` file.
- C: By default, dbt expects model files to be located in a subdirectory named after the target database. You can change this by updating the `model-paths` configuration in your `dbt_project.yml` file.



D: By default, dbt expects model files to be located in a subdirectory named after the target schema. You can change this by updating the model-paths configuration in your dbt_project.yml file.

101. Which of the following best describes configurations in dbt, and how are they applied?
- A:** Configurations are the settings for your target database, and are set in your dbt_project.yml file. They are applied globally across all models.
- B:** Configurations are the SQL statements used to build your models, and are set in your model files. They are applied to individual models.
- C:** Configurations are the parameters used to connect to your target database, and are set in your dbt_project.yml file. They are applied globally across all models.
- D:** Configurations are "model settings" that can be set in your dbt_project.yml file, and in your model file using a config block. They are applied hierarchically, with more specific configurations overriding more general configurations.
- D:** Configurations are "model settings" that can be set in your dbt_project.yml file, and in your model file using a config block. They are applied hierarchically, with more specific configurations overriding more general configurations.
- Explanation:** Configurations in dbt are "model settings" that can be set in your dbt_project.yml file, and in your model file using a config block. Configurations can be used to modify the behavior of a model, such as changing the view name or the target schema. It is important to note that configurations are applied hierarchically, with more specific configurations overriding more general configurations.
102. You are using dbt Core or dbt Cloud to run your project, and you want to ensure that all of your tests are passing. What command should you use?
- C:** dbt test
- The dbt test command is used to execute the tests you defined for your project. This ensures that all of your tests are passing.



- A: dbt run
- B: dbt build
- C: dbt test
- D: dbt deploy

The dbt run command is used to run the models you defined in your project, while the dbt build command is used to build and test resources such as models, seeds, snapshots, and tests. The dbt deploy command is not a valid command in dbt.

103. You are a data analyst working for a subscription-based e-commerce company. The company has recently acquired a new set of data for their customer management system. The data includes information about customers, their subscriptions, and their payment methods. The data is stored in a SQL database and you will be using a tool called dbt to perform data modeling and testing. You have been tasked with ensuring the data from the new set meets the company's requirements for accuracy, completeness, and compliance with financial regulations. Choose the appropriate core dbt tests that should be performed using dbt on the data from the new set, to ensure the data meets the company's requirements:

- A: Test for unique values in the customer identification numbers.
- B: Test for null values in the subscription data.
- C: Test for accepted values in the payment method data, ensuring that the payment methods are only one of the three options

A, B, C, D & E



provided by the company (e.g. credit card, debit card, PayPal)

D: Test for referential integrity between the customer and subscription tables, ensuring that every subscription has a corresponding customer.

E: Test for unique values in the subscription identification numbers.

F: Test that the data is compliant with financial regulations, such as ensuring that credit card numbers are properly encrypted and not accessible to unauthorized parties.

G: Test for the validity of the billing address data by cross-referencing with a geolocation database

104. A data analytics company wants to track changes in their customer database, including hard deletes, to ensure that their analytics are up-to-date. They want to use dbt snapshots to achieve this. What should they do to enable tracking of hard deletes?

A: Turn on the "invalidate_hard_deletes" option in the dbt configuration file.

B: Change the "updated_at" column to a boolean value.

C: Change the snapshot strategy to "check".

D: None of the above.

A: Turn on the "invalidate_hard_deletes" option in the dbt configuration file.

To enable tracking of hard deletes, the "invalidate_hard_deletes" option should be turned on in the dbt configuration file. This is done by setting "invalidate_hard_deletes: true". This will cause dbt to mark rows as invalid if they are deleted from the source query.

105.



What is the main difference between the way snapshots and models are stored in dbt?

A: Snapshots are stored in a separate schema for each user, while models are stored in the same target_schema for everyone.

B: Snapshots are stored in a separate schema for each project, while models are stored in a separate schema for each user.

C: Snapshots are stored in the same target_schema for everyone, while models are stored in a separate schema for each user to maintain separate development and production environments.

D: Snapshots and models are stored in the same schema, but snapshots are stored as binary data while models are stored as SQL files.

C: Snapshots are stored in the same target_schema for everyone, while models are stored in a separate schema for each user to maintain separate development and production environments.

The reason for this difference is that snapshots are meant to be run regularly, and it's best to reference the production version of the snapshot, even when developing models, to ensure consistency and make it easier to build models.

106. You are working on a data warehouse and want to build dependencies between two models that have the same name. In order to achieve this, which of the following options should you use?

A: The ref function with the model name as the argument

B: Custom schemas and custom aliases

C: The ref function with the model name and the schema name as arguments

D: The ref function with the model name and the alias name as arguments

B: Custom schemas and custom aliases
To build dependencies between models, you need to use the ref function. The ref function only takes one argument — The model name (I.E. the filename). As a result, these model names need to be unique, even if they are in distinct folders. Often, this question comes up because users want to give two models the same name in their warehouse, splitting them across separate schemas (E.G. stripe.users and



app.users), to achieve this the best way is to use custom schemas and custom aliases.

107. **You are working on a dbt project and want to change the default location of your seed files. Which of the following best describes the recommended way to implement this?**
- A:** Update the seed-paths configuration in your dbt_project.yml file with the new location of your seed files.
- B:** Create a new subdirectory within the seeds directory and move your seed files to this directory.
- C:** Create a new directory outside of the seeds directory and move your seed files to this directory.
- D:** It is not possible to change the default location of your seed files in dbt.
- A:** Update the seed-paths configuration in your dbt_project.yml file with the new location of your seed files.
- Explanation:** According to dbt documentation, to change the default location of your seed files, you should update the seed-paths configuration in your dbt_project.yml file with the new location of your seed files. For example, if your seed files are located in a directory named "custom_seeds_directory", you can update the configuration like so: seed-paths: ["custom_seeds_directory"]
108. **A data engineer is working on a dbt project and wants to change the relation identifier for a specific model in the project. They know that dbt uses the model file name as the default relation identifier, but they want to use a different identifier for this model. They have read about the alias model configuration and want to know more about it. What is the alias model configuration in dbt?**
- A:** It is a configuration that allows dbt to create relations in a database with a different identifier than the model file name.



B: It is a configuration that allows dbt to automatically create aliases for models in the project based on the table or view name in the database.

C: It is a configuration that allows dbt to create views in the database with a different identifier than the model file name.

D: It is a configuration that allows dbt to automatically create aliases for models in the project based on the schema name in the database.

109. Emily is working on a dbt project and receives an error message indicating a syntax error in the metric definition when trying to run the dbt run command. What is the most likely reason for this issue?

A: The metric definition is missing necessary dimensions.

B: The metric definition contains invalid SQL syntax.

C: The table being used to calculate the metric is not properly defined or does not exist.

D: Too many dimensions are being used to accurately represent the data.

110. While working on a dbt project, Roger needs to provide a comprehensive description for a model that includes multiple paragraphs and formatting. What is the best way to achieve this in dbt?

B: The metric definition contains invalid SQL syntax.

D: Write the description in a .md file and use a doc blocks macro to include it in the schema.yml file.



A: Use a single line description without any formatting.

B: Split the description over multiple lines using "|" and use Markdown.

C: Use the ">" flag to remove interior line breaks and avoid Markdown.

D: Write the description in a .md file and use a doc blocks macro to include it in the schema.yml file.

E: None of the above

111. You are working on a dbt project that has a model that references a source table. You need to modify the source table, and you want to make sure that the model will be updated to reflect the changes. What should you do?

A: Delete the model and recreate it.

B: Run the dbt seed command to reload the source data.

C: Run the dbt run command to rebuild the model.

D: Run the dbt test command to verify the model.

C: Run the dbt run command to rebuild the model.

112. You are working on a dbt Python model and need to understand the behavior of DataFrame operations. Which of the following statements is true regarding DataFrame operations in dbt Python models?

A: All operations are immediately comput-

B: DataFrame operations are lazily evaluated, which means that the results are not immediately calculated, and the final result is only computed when explicitly asked for.



ed, and the final result is saved as a table in your data warehouse.

B: DataFrame operations are lazily evaluated, which means that the results are not immediately calculated, and the final result is only computed when explicitly asked for.

C: DataFrame operations are cached by dbt, which means that the final result is always available even without explicitly asking for it.

D: The preview of data in development is not possible with DataFrame operations in dbt Python models.

113. Which of the following is true about dbt seeds?

A: Seeds can be used to provide reference data that does not change frequently

B: Seeds can be used to provide initial data for a new project

C: Seeds can only be loaded into a PostgreSQL database

D: Seeds can be updated through SQL queries

A & B: Seeds can be used to provide reference data that does not change frequently,

Seeds can be used to provide initial data for a new project

Explanation: dbt seeds can be used to provide reference data that does not change frequently, such as a list of countries or product categories. They can also be used to provide initial data for a new project.

114. You are setting up a snapshot configuration in dbt and want to ensure that the unique key is reliable. What is a best practice for ensuring the uniqueness of the unique key?

A: Add a test for uniqueness to the snapshot table.

B: Add a test for uniqueness to the source data.



B: Add a test for uniqueness to the source data.

C: Add a test for uniqueness to the target data.

D: None of the above

115. You are using dbt to connect to a warehouse. You want to check if the connection is successful. What command should you use?

A: dbt debug

Run dbt debug from within a dbt project to test your connection

A: dbt debug

B: dbt test

C: dbt run

D: dbt check

116. You are a data analyst working for a large retail company. The company has recently acquired data from a new supplier for their inventory management system. The data includes information about products, suppliers, and orders. The data is stored in a SQL database. You have been tasked with ensuring the data from the new supplier meets the company's requirements for accuracy and completeness. Choose the appropriate dbt generic tests and assumptions that should be made and verified for the data from the new supplier:

A, B, C, D & E: The product names in the data are unique and do not contain any null values, The supplier names in the data are not null and have a maximum length of 100 characters, The order numbers in the data are unique and have a valid format (e.g. "ORD-####"), There is a referential integrity between the product and supplier tables, meaning that every product in the product table has a corresponding supplier in the supplier table, The prices in the data have accepted values within a specified range (e.g. £0.01- £1000)

A: The product names in the data are unique and do not contain any null values

B: The supplier names in the data are not null and have a maximum length of 100



characters

C: The order numbers in the data are unique and have a valid format (e.g. "ORD-####")

D: There is a referential integrity between the product and supplier tables, meaning that every product in the product table has a corresponding supplier in the supplier table

E: The prices in the data have accepted values within a specified range (e.g. £0.01-£1000)

F: All fields in the data are not null

-
117. In dbt, what is the recommended approach to create default values for a column in a model?

A: Use the default configuration block in the model file to specify a default value.

B: Use the `set_default` function to set a default value in the SQL select statement.

C: Use DDL to create a default value for the column in the database.

D: Use SQL to express the default value in the SQL select statement.

D: Use SQL to express the default value in the SQL select statement.

-
118. In a dbt project, after completing the data modeling process and testing the model to ensure everything is working correctly, which of the following steps is crucial for deploying the project to production and keeping the data up-to-date?

B: Scheduling a job to run the code periodically



- A: Rebuilding the model manually every time new data is added
 - B: Scheduling a job to run the code periodically
 - C: Converting the model's materialization to a view
 - D: Redefining the model's unique key
-

119. Which of the following statements accurately describe the dbt ref macros?
- A: The dbt ref function allows users to select from existing tables and views in their warehouse.
 - B: The dbt ref function allows users to implement transformations in stages by selecting from another model.
 - C: The dbt ref function allows users to create new tables and views in their warehouse.
 - D: The dbt ref function does not allow users to specify the order in which models are executed.
120. What is the purpose of a dbt metric?
- A: To define the raw data used in dbt transformations.
 - B: To describe the downstream use of a dbt project.
 - C: To monitor the performance of dbt transformations.
 - D: To define important metrics for a dbt project.



121. Maria is a data engineer working with dbt for a client. She has identified a need to execute custom SQL specific to the client's data platform that is not provided by dbt's built-in functionality. What feature should Maria use to accomplish this?
- B: Hooks
- Explanation: Use hooks to execute custom SQL specific to the client's data platform that is not provided by dbt's built-in functionality. Hooks can also call macros that return SQL statements.
- A: Seeds
B: Hooks
C: Macros
D: Packages
-
122. When using dbt CLI, which command should be run first to generate necessary documentation artifacts before serving the model to a local web server?
- A: dbt docs generate
B: dbt docs serve
C: dbt deploy
D: dbt cloud
- A: dbt docs generate
-
123. In the freshness block property of the source schema.yml file, which properties control the behavior of the source and what happens if neither is provided?
- C: warn_after and error_after control the behavior of the source; if neither is provided, dbt will not calculate freshness snapshots for the tables in the source.
- A: error_before and max_age control the behavior of the source; if neither is provided, dbt will not calculate freshness snapshots for the tables in the source.
B: max_age and warn_before control the behavior of the source; if neither is provided, dbt will not calculate freshness snapshots for the tables in the source.



C: warn_after and error_after control the behavior of the source; if neither is provided, dbt will not calculate freshness snapshots for the tables in the source.

D: error_after and warn_before control the behavior of the source; if neither is provided, dbt will not calculate freshness snapshots for the tables in the source.

124. You have created a Python model in dbt and ran "dbt run --select python_model.py". What does this command do?

A: It will run all the models in the project.

B: It will only run the models with a specific tag.

C: It will only run the Python model.

D: It will only run the models with a specific schema.

C: It will only run the Python model.

125. You have an incremental model in dbt with the default setting "on_schema_change: ignore", but when running the dbt command, the removed column is still present in the target table. What could be the cause of this issue?

A: The on_schema_change parameter is not set correctly in the configuration.

B: The database does not support removing columns in this way.

C: A full-refresh of both the incremental model and any related models has not been executed.

C: A full-refresh of both the incremental model and any related models has not been executed.

Explanation: When using an incremental model in dbt with the "on_schema_change: ignore" setting, any changes made to the columns in the incremental model will not be reflected in the target table until a full-refresh of both the incremental model and any related models is executed. This means that if you remove a column from the incremental model, it will still remain in the target table until a full-re-



D: There is an issue with the SQL used in the model.

fresh is executed. Therefore, the most likely cause of this issue is that a full-refresh of both the incremental model and any related models has not been executed.

126. You are working on a dbt project and have defined several metrics. Before deploying the project, you want to preview and verify the metrics within dbt. Which of the following steps should you take to achieve this?

A: Manually run the SQL code for each metric in the data warehouse

B: Add the dbt_metrics package installation config to your packages.yml file, run dbt deps to install the package, and then run the metrics.calculate macro with the defined metrics

C: Change the materialization of the models containing the metrics to "view"

D: Modify the metrics definition in the schema.yml file to include a preview flag

B: Add the dbt_metrics package installation config to your packages.yml file, run dbt deps to install the package, and then run the metrics.calculate macro with the defined metrics

Explanation: To preview and verify metrics within dbt before deploying, you need to have the dbt_metrics package installed. This can be done by adding the metrics package installation code to your packages.yml file and running dbt deps to install the package. Once the package is installed, you can use the metrics.calculate macro with the defined metrics to preview and verify them.

127. You are reviewing a dbt project that has a source table defined in the schema.yml file. The table has a name in the database that looks like machine generated name which is not easily readable, you want to use a more sensible name in dbt. How can you implement this and still ensure dbt compiles using the actual table name?

A: Use DDL and DML statements to change the table name to one that can be easily

C: Use the identifier property to define the actual table name as per the database, and use your name: property for the name that makes sense and is easily readable



read and makes sense

B: Use the 'Table' property to define the actual names as per the database, and use your name: property for the name that makes sense and is easily readable

C: Use the identifier property to define the actual table name as per the database, and use your name: property for the name that makes sense and is easily readable

D: Use the 'Model' property to define the actual names as per the database, and use your name: property for the name that makes sense and is easily readable

-
128. Which of the following is NOT a benefit of using a code templating tool like Jinja in dbt? C: It allows for the creation of new database tables

A: Control structures can be used in queries

B: Macros can be created to reuse SQL code

C: It allows for the creation of new database tables

D: It helps to optimize the performance of queries

-
129. You are working on a dbt project that involves organizing SQL queries in a directory called "reports" without materializing them to your data platform. You want to define the location of the reports directory in the dbt_project.yml file. What configura- C: analysis-paths
- Explanation: To define the location of the reports directory in the dbt_project.yml file, you should use the analysis-paths configuration option. This option specifies the location of the directory where analytical



tion option should you use?

A: models-paths

B: seeds-paths

C: analysis-paths

D: macro-paths

SQL queries are stored. Option A is incorrect because the models configuration option is used to specify the location of the directory where dbt models are stored. Option B is incorrect because the seeds configuration option is used to specify the location of the directory where dbt seed data is stored. Option D is incorrect because the macro-paths configuration option is used to specify the location of the directory where dbt macros are stored.

130. You are setting up a new dbt project for your organization and need to configure the connection profile. Which of the following steps should you take to ensure the correct setup?

A: Create a new profile in your data warehouse with the name of your organization and grant the necessary permissions

B: Copy the appropriate sample profile for your warehouse into your profiles.yml file, update the details with your organization's name, and ensure it matches the profile indicated in your dbt_project.yml file

C: Modify the dbt_project.yml file to include your organization's name and the warehouse connection details

D: Add a new environment variable with your organization's name and the warehouse connection details

B: Copy the appropriate sample profile for your warehouse into your profiles.yml file, update the details with your organization's name, and ensure it matches the profile indicated in your dbt_project.yml file

Explanation: To set up your profile, copy the correct sample profile for your warehouse into your profiles.yml file and update the details as follows: Profile Name: Replace the name of the profile with a sensible name - it's often a good idea to use the name of your organization. Make sure that this is the same name as the profile indicated in your dbt_project.yml file.



131. Your team is developing a new data warehouse and is deciding on which tools to use. You are considering using dbt as your Transformation tool. What are the benefits of using dbt as your Transformation tool?
- A: dbt is an open-source tool that is easy to learn and use, making it a cost-effective solution.
- B: dbt integrates easily with cloud-based data warehouses, making it a scalable solution for handling large amounts of data.
- C: dbt provides a declarative language for defining your data transformations, which makes it easier to maintain and troubleshoot your code.
- D: All of the above.
-
132. You are working on a dbt project and are considering using seeds to provide reference data for your models. Which of the following is a good use-case for seeds according to the best practices in dbt?
- A: Loading raw data that has been exported to CSVs
- B: A list of mappings of country codes to country names
- C: Any kind of production data containing sensitive information such as PII and passwords
- D: A list of log files containing user activity on a website
- B: A list of mappings of country codes to country names
- Explanation: According to the best practices in dbt, a good use-case for seeds is to provide reference data for your models, such as lookup tables or configuration data. A list of mappings of country codes to country names is a good example of this. Other examples might include a list of test emails to exclude from analysis or a list of employee account IDs.
-



133. **Minta is working on a dbt project and has defined both pre- and post-hooks for a model in the configuration block of the model file. She later realizes that she needs to add an additional post-hook to the model. Which of the following is true?**
- A: She can only define additional hooks in the configuration block of the model file**
 - B: She can only define additional hooks in the dbt_project.yml file**
 - C: She can define additional hooks in both the configuration block of the model file and the dbt_project.yml file**
 - D: She cannot add additional hooks to the model once it has been defined**
134. **What is the purpose of using Jinja statements in dbt projects?**
- A: To store configuration data for use in dbt projects**
 - B: To generate reports and visualizations from SQL queries**
 - C: To generate SQL code dynamically based on configuration data or user input**
 - D: To define data transformations in a graphical user interface**
- C:** She can define additional hooks in both the configuration block of the model file and the dbt_project.yml file
- C:** To generate SQL code dynamically based on configuration data or user input
- Explanation:** In dbt projects, Jinja statements are primarily used to generate SQL code dynamically based on configuration data or user input. This can be useful in a variety of scenarios, such as generating queries that join multiple tables, filtering data based on user input, or creating dynamic reports. By using Jinja statements for control flow, you can make your dbt projects more flexible and adaptable to changing requirements. Options B, C, and D are not correct because Jinja statements are not typically used to store configuration data, generate reports, or define data



transformations in a graphical user interface.

-
135. **A data team is creating a snapshot in dbt to track changes to their customer data over time. They are unsure of whether they should include all columns in the snapshot, even if some columns do not seem useful. Which of the following is a best practice for selecting columns for a snapshot in dbt?**
- A: Only include the columns that seem useful at the moment
- B: Include as many columns as possible in the snapshot
- C: Include only the columns that are necessary for downstream models
- D: Include only the columns that are unique to the table being snapped
- B: Include as many columns as possible in the snapshot
-
136. **Which function do you use to build dependencies between models in dbt?**
- A: The ref function
- B: The unique function
- C: The dependency function
- D: The model function
- A: The ref function
-
137. **What is the requirement for a custom calendar table in dbt metrics?**
- A: It must contain a "date_day" column and columns for week, month, and quarter, and year (or equivalent) with the prefix "date_".
- B: It must contain a "date_day" column and
- A: It must contain a "date_day" column and columns for week, month, and quarter, and year (or equivalent) with the prefix "date_".
- Explanation: To use a custom calendar table in dbt metrics, it must contain a



columns for week, month, and quarter with the prefix "date_" only.

C: It must contain a "date_day" column and columns for week, month, quarter, and year (or equivalent) with any prefix.

D: It must contain a "date_day" column only.

"date_day" column and columns for week, month, and quarter, and year (or equivalent) with the prefix "date_". Any additional date columns must also have the "date_" prefix. This is because the dbt metrics package includes a basic calendar table for your dbt runs, and to use a custom calendar, it must follow this format.

138. A data team wants to implement snapshot tables in their dbt project to track changes to a large dataset. What is the best practice for working with snapshots in dbt?

A: Store snapshots in the same schema as other models.

B: Store snapshots in a separate schema and set different privileges for them compared to other models.

C: Store snapshots in the same schema but with a different file extension than other models.

D: A and B

B: Store snapshots in a separate schema and set different privileges for them compared to other models.

139. What happens if a syntax error is present in one of the SQL expressions specified in the "incremental_predicates" feature?

A: The feature will be disabled.

B: The syntax error will be ignored and the feature will work as expected.

C: The syntax error will cause an error and prevent the feature from working.

C: The syntax error will cause an error and prevent the feature from working.



D: The syntax error will cause data corruption.

140. **What is the best practice for managing user credentials in dbt?**

A: Assign a single set of database credentials for all dbt users.

B: Assign a separate set of database credentials for each dbt user, including a separate user for production runs.

C: Assign a single set of database credentials for all dbt users, but have different privileges for different users.

D: Assign a single set of database credentials for all dbt users, but have different privileges for different data warehouses.

B: Assign a separate set of database credentials for each dbt user, including a separate user for production runs.

This helps debug rogue queries, simplifies ownership of schemas, and improves security.

141. **A data team is implementing dbt snapshots to track changes in their customer and order tables. They want to ensure they have a unique key for each record to match and track changes over time. Which column in a snapshot is used by dbt to match and track records as they change over time?**

A: dbt_valid_from column

B: dbt_valid_to column

C: dbt_scd_id column

D: updated_at column

C: dbt_scd_id column

dbt_scd_id is a unique key generated for each snapshot record. This is used internally by dbt

142. **A company is using dbt to track changes in their customer data. They have specified**

D: No action is taken by dbt.

Explanation: If the values in the specified



the `check_cols` configuration in their snapshot block. What happens if the values in the specified columns remain unchanged?

A: The old data in that row is considered outdated and replaced with the new data.

B: The `updated_at` value in the row is compared to the last time dbt checked.

C: Dbt recognizes that the data in the row has been altered.

D: No action is taken by dbt.

columns remain unchanged, dbt will not take any action. The check strategy only considers the old data in a row outdated and replaces it with the new data if any of the specified columns have different values than their historical values.

143. A data analytics company wants to keep track of changes in their sales data, specifically in the orders table. They want to see how the total order value and the number of items per order have changed over time. Which of the following statements about using snapshots in dbt is correct?
- A: Snapshots are not suitable for tracking changes in sales data as this requires aggregating data over time, which is not possible with snapshots.
- B: To use snapshots, the company needs to write a select statement that retrieves the desired data and wrap it in a snapshot block in a .sql file located in the "snapshots" directory.
- C: Snapshots are only used to track changes in dimension tables, such as changes to a customer's address or name, over time.
- D: The company can use snapshots by in-
- B: To use snapshots, the company needs to write a select statement that retrieves the desired data and wrap it in a snapshot block in a .sql file located in the "snapshots" directory.
- While snapshots are typically used for tracking changes in dimension tables, they can also be used for tracking changes in fact tables like the orders table by selecting the relevant columns and aggregating them over time. The unique key is also an important part of configuring snapshots, but it is not the only required component.



cluding a "unique key" in their snapshot configuration, which allows dbt to identify changes in records by comparing the key value to previous snapshots.

144. As a data analyst working with dbt Cloud, how can you check the SQL that dbt is running for a specific model?

Choose only ONE best answer.

A: Check the dbt Cloud homepage for a list of recently run models, and click on the model name to view the SQL.

B: Look for a "SQL" button on the toolbar in the dbt Cloud web interface, and click it to view the SQL for the currently selected model.

C: Check within the run output, click on a model name, and then select "Details".

D: Run the "dbt debug" command on your local machine to view the SQL for a specific model.

145. What is the recommended approach to break down complex SQL code when using dbt data models?

A: Breaking complex CTEs into separate data models

B: Combining all logic into one large data model

C: Utilizing only intermediate models

D: Using only the final model to produce the desired result set



-
146. You are attempting to define a new metric in your dbt project using an ephemeral model as the source materialized data. When you attempt to run the dbt run command, you receive an error message that says "Error: Cannot create metric from temporary data model." What is the most likely reason for this issue?
- A: The metric definition is not properly defined in the YAML file.
- B: The metric name contains invalid characters or exceeds the maximum length.
- C: Ephemeral models cannot be used as a source for defining metrics.
- D: The dbt_project.yml file is missing or improperly configured.
- C: Ephemeral models cannot be used as a source for defining metrics. Ephemeral models are temporary data models without a permanent representation in the data warehouse, which is necessary when defining metrics.
-
147. What is the main advantage of using common table expressions (CTEs) in SQL?
- A: Ensures modularity
- B: Makes the code easier to read and maintain
- C: Separates logic that cleans up data from logic that transforms data
- D: B and C
- E: None of the above
- D: B and C
-
148. Your dbt Cloud deployment has failed and you need to troubleshoot the issue. Which of the following steps should be part of your troubleshooting process?
- A: Check the log files on the server for any
- D: All of the above



errors

B: Reach out to the dbt support team for assistance and provide them with the deployment script

C: Set up email and Slack notifications for any failed runs in the Account Settings > Notifications and review the deployment script for any syntax errors or logical mistakes

D: All of the above

149. What is the recommended location for the profiles.yml file that contains the details needed to connect to the data warehouse using dbt?
- A: Inside the dbt project
 - B: In version control
 - C: Outside the dbt project
 - D: In the dbt_project.yml file
- C: Outside the dbt project
This file is generally located outside of the dbt project to avoid sensitive credentials being checked into version control. However, it can be safely checked in when using environment variables to load sensitive credentials. The profiles.yml file is searched for in the current working directory, and if not found, dbt defaults to looking in the ~/.dbt/ directory.
-
150. You are running a model in dbt and the SQL code contains an error. dbt returns the error message from the database or data warehouse and skips any models downstream of the affected model. What should you do to debug the error?
- A: Create a new model and try running the SQL code again
 - B: Use the error message and the compiled SQL to troubleshoot the error



C: Check the data in the database or data warehouse for any issues

D: Ignore the error and move on to the next model

151. Your team has created a schema.yml file for a dbt project. One of your team members is unclear about how to define sources in the file. How would you explain the process to them?

A: In the schema.yml file, sources are defined under the tables: key. Each source is defined with a name, and a list of tables associated with that source.

B: In the schema.yml file, sources are defined under the models: key. Each source is defined with a name, and a list of tables associated with that source.

C: In the schema.yml file, sources are defined under the sources: key. Each source is defined with a name, and a list of tables associated with that source.

D: In the schema.yml file, sources are not defined directly. They are automatically detected based on the tables included in the models.

C: In the schema.yml file, sources are defined under the sources: key. Each source is defined with a name, and a list of tables associated with that source.

152. In dbt metrics, if you want to use a custom calendar table, what column must the table contain?

A: week, "month", "quarter", and "year" columns

B: date_day column and columns for week, month, quarter, and year (or equivalent) with the prefix "date_"

Explanation: In dbt metrics, if you want to use a custom calendar table, the table



B: date_day column and columns for week, month, quarter, and year (or equivalent) with the prefix "date_"

C: date_day, "week", "month", and "year" columns

D: date_day column and a "custom_calendar" column

must contain a "date_day" column and columns for week, month, quarter, and year (or equivalent) with the prefix "date_". Any additional date columns must also have the "date_" prefix. This allows dbt metrics to recognize the date-related columns and use them in calculations. To specify the custom calendar, you need to change the value of the "dbt_metrics_calendar_model" variable in the "dbt_project.yml" file.

153. A dbt project contains multiple models that depend on a snapshot. How can the analytics engineers ensure that the downstream models always use the most up-to-date snapshot data?
- A: By using the "model" function in queries
- B: By manually updating the downstream models every time the snapshot changes
- C: By specifying the snapshot as a source in the downstream models
- D: By creating a separate model to act as an intermediary between the snapshot and the downstream models
- C: By specifying the snapshot as a source in the downstream models
154. Which of the following is true about snapshots in dbt?
- Chose ALL answers that apply:
- A: Snapshots can be used to version control tables
- B: Snapshots can be used to capture the state of mutable tables
- B & D: Snapshots can be used to capture the state of mutable tables, Snapshots can be used to compare changes between versions of a table



C: Snapshots are only applicable to certain types of databases

D: Snapshots can be used to compare changes between versions of a table

155. You are a data analyst working on a dbt project. You want to use the built-in generic tests that come with dbt to enhance the reliability of your project. Which of the following is not a built-in generic test in dbt?
- A: unique
B: not_null
C: all_anomalous
D: accepted_values
E: relationships

156. A dbt project contains a snapshot with a unique key that is not truly unique. What is the potential consequence of this?
- A: The snapshot will be automatically deleted by dbt
B: The snapshot will become corrupt and unusable
C: dbt will not be able to accurately track changes to the data over time
D: The snapshot will not be accessible to downstream models
- C: dbt will not be able to accurately track changes to the data over time
The unique key in the snapshot configuration is used by dbt to match rows in the source data with the corresponding rows in the snapshot table. If the unique key is not unique, dbt will not be able to accurately track changes to the data over time.

157. Where should you store the profile.yml file when working with dbt?
- A: In the root directory of your project
B: in the dbt_modules folder
C: in a config folder outside of your dbt project



C: in a config folder outside of your dbt project
D: on the local machine

158. What does the dbt audit_helper package do?

Chose ALL answers that apply

A: Writes your own query manually

B: Generates comparison queries between before and after states

C: Compares original query results to refactored results

D: Identifies variances more quickly

B & C & D: Generates comparison queries between before and after states, Compares original query results to refactored results, Identifies variances more quickly

159. You are working on a dbt project that involves working with multiple packages to process and transform data. You want to ensure that the specific python package versions you are using are tracked for better implementation and tracking purposes. Which of the following is the recommended way to specify package versions in the dbt project metadata?

A: Include a separate requirements.txt file in the dbt project directory that lists all the necessary packages and their versions.

B: Specify the package versions in the dbt_project.yml file under the "packages" configuration block.

C: Specify the package versions in each individual model file that uses the packages.

D: Include the package versions in the SQL

B: Specify the package versions in the dbt_project.yml file under the "packages" configuration block.



code used to manipulate and transform the data.

160. You are working on a dbt transformation model that uses data from a data source. You want to describe the downstream use of the project. What should you do?
- A: Add a comment to the relevant SQL code.
- B: Include the source in the dbt project README.
- C: Create a dbt exposure for the project.
- D: Add a note to the dbt source description.
- C: Create a dbt exposure for the project. The purpose of a dbt exposure is to describe the downstream use of a dbt project. It provides a better overview of what the project is being used for, helping to inform decision-making and support collaboration between teams
-